



TreScout

DAILY TECHNOLOGY REPORT

Today's Technology Trends

August 4, 2026 (Tuesday)

A FEW SIPS

Today, discussions that developer tools should be open source and new architectures that reward the level of expertise of high-performance language models (LLMs) come to the fore. Data on GitHub and HuggingFace show that efforts to run large models with low memory consumption (inference) and 3D data compression techniques are gaining momentum.

34 HIGHLIGHTS · ~51 MIN READ

16 GitHub
Trending

5 Hacker
News

5 HuggingFace ·
Daily Models

5 HuggingFace ·
Daily Papers

3 Lobsters

01 lyogavin/airllm

AirLLM allows you to run large language models with 70 billion parameters on graphics processors (GPU) with only 4 gigabytes of video memory. This library increases accessibility by running models with high hardware requirements on low-capacity consumer-grade hardware.

[Jupyter Notebook](#) · ★ 27,515 · +1,085 today <https://github.com/lyogavin/airllm>

02 zhaoxuya520/reverse-skill

Reverse-skill is an AI-powered talent driver that automates reverse engineering and penetration testing processes. By working integrated with artificial intelligence-supported coding tools (AI coding clients) such as Claude Code or Cursor, it automatically establishes the tool chains needed for security research and constantly updates its experience base.

[PowerShell](#) · ★ 16,395 · +2,446 today <https://github.com/zhaoxuya520/reverse-skill>

03 firecrawl/pdf-inspector

PDF Inspector, developed by Firecrawl, is a fast Rust library designed to examine, classify and extract the text of PDF files. It enables intelligent routing decisions to be made in data processing processes by distinguishing scanned documents from text-based files.

[Rust](#) · ★ 8,766 · +1,699 today <https://github.com/firecrawl/pdf-inspector>

04 esengine/DeepSeek-Reasonix

DeepSeek-Reasonix is an AI coding agent that runs on the terminal and integrates with DeepSeek models. Focusing on prefix-cache stability, this tool allows developers to manage long-term coding processes without interruption.

[Go](#) · ★ 30,213 · +883 today <https://github.com/esengine/DeepSeek-Reasonix>

05 TencentCloud/TencentDB-Agent-Memory

TencentDB Agent Memory is a hub for AI agents that transforms conversations, documents, and code into reusable memory assets. This system makes data shareable between different agents by structuring it as chat history, skill, knowledge base (LLM-Wiki) and Code-Graph.

[TypeScript](#) · ★ 12,417 · +1,090 today <https://github.com/TencentCloud/TencentDB-Agent-Memory>

06 microsoft/AI-For-Beginners

AI-For-Beginners, prepared by Microsoft, offers a 12-week curriculum for those who want to enter the world of artificial intelligence. This training set in Jupyter Notebook format, containing interactive code blocks, provides a wide learning path from basic concepts to applied projects.

[Jupyter Notebook](#) · ★ 61,100 · +1,902 today <https://github.com/microsoft/AI-For-Beginners>

07 microsoft/generative-ai-for-beginners

Prepared by Microsoft, generative-ai-for-beginners offers a 21-lesson training curriculum for those who want to get into generative artificial intelligence (generative AI) technologies. The content aims to teach basic concepts and model development processes through applied projects in Jupyter Notebook format.

Jupyter Notebook · ★ 115,799 · +775 today <https://github.com/microsoft/generative-ai-for-beginners>

08 donnemartin/system-design-primer

System Design Primer offers a comprehensive learning resource for those who want to design large-scale software systems. It brings together system design principles and practical study cards for software developers preparing for technical interviews.

Python · ★ 360,710 · +237 today <https://github.com/donnemartin/system-design-primer>

09 antirez/ds4

Developed by the creator of Redis, ds4 is an inference engine that allows running DeepSeek models on local devices. With CUDA and ROCm support, Metal offers hardware-accelerated artificial intelligence models optimized with low-level C language.

C · ★ 20,440 · +384 today <https://github.com/antirez/ds4>

10 shiyu-coder/Kronos

Kronos is a foundation model designed to understand and analyze the complex language of financial markets. Trained on financial data, this Python-based system aims to automate the processes of interpreting market movements and processing financial texts.

Python · ★ 35,882 · +200 today <https://github.com/shiyu-coder/Kronos>

11 Panniantong/Agent-Reach

Agent-Reach is a Python tool that allows AI agents to search and read content on various platforms on the internet (Twitter, Reddit, YouTube, GitHub). It allows access to web data via the command line without paying any application programming interface (API) fee.

Python · ★ 65,979 · +1,057 today <https://github.com/Panniantong/Agent-Reach>

12 Alishahryar1/free-claude-code

Developed by Alishahryar1, free-claude-code provides free access to artificial intelligence models such as Claude Code and Codex via terminal, application or development environment (IDE). This open source tool integrates with platforms that offer voice command support, such as OpenClaw, allowing users to use artificial intelligence capabilities on different devices.

Python · ★ 44,150 · +278 today <https://github.com/Alishahryar1/free-claude-code>

13 **iv-org/invidious**

Invidious is an alternative interface developed for the video platform YouTube. It offers users a viewing experience that is ad-free and does not track personal data.

Crystal · ★ 22,316 · +402 today <https://github.com/iv-org/invidious>

14 **livekit/agents**

LiveKit Agents is a Python framework used to develop real-time audio and video artificial intelligence agents. It enables developers to create low-latency voice interactions, making it easier to create AI systems that respond human-like.

Python · ★ 12,122 · +148 today <https://github.com/livekit/agents>

15 **usekaneo/kaneo**

Kaneo is an open source project management tool that delivers a user-centric experience by streamlining complex features. Developed with TypeScript, this platform aims to create a simplified alternative to popular project tracking software such as Jira.

TypeScript · ★ 6,988 · +665 today <https://github.com/usekaneo/kaneo>

16 **jamiepine/voicebox**

Voicebox is an open source artificial intelligence voice studio that allows users to perform voice cloning, dictation and voice production. Developed with TypeScript language, this tool offers a comprehensive solution for those who want to run voice technologies in a local environment.

TypeScript · ★ 48,852 · +412 today <https://github.com/jamiepine/voicebox>

01 LLMs reward expertise

Large language models (LLM) show higher performance on complex tasks when provided with specialized contexts. Increasing the technical knowledge of users directly affects the quality of the outputs received from artificial intelligence models.

717 points · 300 comments <https://www.seangoedecke.com/llms-reward-expertise/>

Discussion: news.ycombinator.com/item?id=49161518 (717 oy · 300 yorum)

02 Devtools must be open source

Open source software development tools (devtools) ensure long-term sustainability of the tools by reducing user dependency. This approach allows developers to have full control over the technologies they use and make community-supported improvements.

562 points · 195 comments <https://blog.exe.dev/devtools-must-be-open-source>

Discussion: news.ycombinator.com/item?id=49156111 (562 oy · 195 yorum)

03 Ten advances in mathematics and theoretical computer science

OpenAI has compiled ten important advances in mathematics and theoretical computer science. This study summarizes new methods used in solving complex problems and current developments in the field of computational mathematics.

510 points · 781 comments <https://openai.com/index/ten-advances-in-mathematics/>

Discussion: news.ycombinator.com/item?id=49157930 (510 oy · 781 yorum)

04 More German than many Germans

Mert Bulan conveys the complex relationship between bureaucratic processes and cultural habits in Germany and digitalization through his own experiences. The article discusses the efforts to adapt the traditional functioning of the country to modern technology and the difficulties encountered in this process.

468 points · 325 comments <https://mertbulan.com/more-german-than-many-germans/>

Discussion: news.ycombinator.com/item?id=49151734 (468 oy · 325 yorum)

05 Prevent cognitive debt by manually retyping LLM-generated code

Manually rewriting the codes produced by artificial intelligence instead of copying and pasting them allows developers to understand the logic more deeply and prevent the formation of cognitive debt. This method supports debugging errors and better ownership of the system by increasing awareness of the functioning of the code.

447 points · 364 comments <https://ankursethi.com/blog/prevent-cognitive-debt-by-manually-retyping-llm-generated-code/>

Discussion: news.ycombinator.com/item?id=49153374 (447 oy · 364 yorum)

01 moonshotai/Kimi-K3

Developed by Moonshot AI, Kimi-K3 is a multimodal artificial intelligence model that can process visual and text data simultaneously. It allows users to convert visual content into text and perform interactive analysis on this data.

♥ 9,884 · 967,622 downloads <https://huggingface.co/moonshotai/Kimi-K3>

02 deepseek-ai/DeepSeek-V4-Flash-0731

DeepSeek-V4-Flash-0731, the text-generation model developed by DeepSeek, offers a structure optimized for high-performance language processing tasks. The model provides developers with an accessible language model alternative through the open-source AI platform HuggingFace.

♥ 2,113 · 236,076 downloads <https://huggingface.co/deepseek-ai/DeepSeek-V4-Flash-0731>

03 MiniMaxAI/MiniMax-H3

MiniMax-H3, developed by MiniMaxAI, is an artificial intelligence model that produces videos using visual and text inputs. This model supports video creation processes by converting complex commands into visual content.

♥ 1,602 · 0 downloads <https://huggingface.co/MiniMaxAI/MiniMax-H3>

04 DavidAU/Qwen3.6-27B-Fable-Fusion-711-Uncensored-Heretic-NM-DAU-NEO-MAX-MTP-GGUF

Qwen3.6-27B-Fable-Fusion, shared on Hugging Face, is an artificial intelligence model that produces text output by processing visual and text data. With its uncensored structure and high number of parameters, it allows users to produce complex content on their local systems.

♥ 1,455 · 1,550,034 downloads <https://huggingface.co/DavidAU/Qwen3.6-27B-Fable-Fusion-711-Uncensored-Heretic-NM-DAU-NEO-MAX-MTP-GGUF>

05 Comfy-Org/MiniMax-H3

MiniMax-H3 is a multi-modal large language model with text and visual processing capabilities. ComfyUI works integrated with the visual editing interface, allowing users to manage complex image creation processes faster.

♥ 477 · 2 downloads <https://huggingface.co/Comfy-Org/MiniMax-H3>

01 3DZip: Spatial-Aware Feature Diversity-Guided Token Compression for 3D Question Answering

3DZip offers a spatial awareness-oriented token compression method to reduce the burden of intensive data processing in three-dimensional vision-language models (3D VLM). This technique preserves the scene geometry, eliminates unnecessary data, and increases computational efficiency in three-dimensional question-answer processes.

5 likes · 2 comments <https://huggingface.co/papers/2608.01185>

02 DAPD: Dual-Anchored Policy Distillation

Dual-anchored policy distillation aims to eliminate the information asymmetry between teacher and student models in the training process of language models. This method increases model performance by preventing the student from developing erroneous behaviors based on privileged data that he cannot access during the inference phase.

1 likes · 1 comments <https://huggingface.co/papers/2608.01735>

03 LongHorizon-Harness: Advancing Long-Horizon Agents for Real-World Tasks

LongHorizon-Harness aims to solve the state tracking and error accumulation problems experienced by large language model-based agents in long-term missions. By redefining complex processes as a task-state management problem, it enables agents to make more consistent decisions in multi-step workflows.

44 likes · 1 comments <https://huggingface.co/papers/2608.01964>

04 DiffusionGemma Technical Report

Developed by Google, DiffusionGemma performs parallel processing in blocks of 256 using discrete diffusion, unlike traditional methods in text production. This approach aims to generate text at high speed by overcoming the sequential word generation bottleneck in large language models.

2 likes · 0 comments <https://huggingface.co/papers/2608.00146>

05 CADENA: Stepwise CAD Reverse Engineering

CADENA is a system that converts three-dimensional network models (3D mesh) into parametric computer-aided design (CAD) programs. By mimicking the way human engineers work, it builds the model step by step, adding features and performing geometric verifications at each stage, rather than building it all at once.

1 likes · 1 comments <https://huggingface.co/papers/2608.00799>

01 Kill the Cookie Banner

Kill the Cookie Banner is an initiative that aims to improve user experience by blocking cookie banners on websites at the browser level. It aims to eliminate notification fatigue caused by legal regulations throughout the European Union with automatic approval or rejection mechanisms.

103 points · 20 comments <https://killthecookiebanner.eu/>

Discussion: lobste.rs/s/afttrcr/kill_cookie_banner (103 oy · 20 yorum)

02 Don't be a meat proxy

It is criticized that software developers disable their own logic and become just intermediaries (proxy) when using artificial intelligence tools. This approach is defined as a risk that reduces code quality and dominance over the system.

95 points · 42 comments <https://gruhn.me/blog/2026-08-03/>

Discussion: lobste.rs/s/hfbqr3/don_t_be_meat_proxy (95 oy · 42 yorum)

03 help wanted

Lake is an open source operating system project that aims to redefine the personal computing experience. The development team is looking for new developers to contribute in key areas such as operating system architecture and user interface design.

143 points · 19 comments <https://lake.computer/blog/help-wanted/>

Discussion: lobste.rs/s/a355zh/help_wanted (143 oy · 19 yorum)

large language models

Artificial intelligence models trained with huge data that can understand human language and produce text.

GPU

Powerful processor unit used to process images on the computer and accelerate artificial intelligence calculations.

AI coding clients

AI-powered tools that help you write code, find errors and complete the project while developing software.

prefix-cache

It is a method that allows artificial intelligence to answer the same questions much faster by keeping the information it has previously processed in its memory.

skill

The ability of an artificial intelligence assistant to specialize in a specific subject and perform user-specific tasks.

LLM-Wiki

A digital information library where artificial intelligence models are fed from external sources to provide more accurate information.

Code-Graph

A visual structure that maps the interconnections of different pieces of code in a software project.

artificial intelligence

Computers gaining the ability to think, learn and solve problems like humans.

generative AI

Artificial intelligence technology that can create new texts, images or sounds using existing data.

system design

The process of planning how parts of a complex software will work with each other and how data will flow.

inference engine

The software layer that enables a trained artificial intelligence model to make predictions on new data given to it.

foundation model

The main model that forms the basis of artificial intelligence applications, trained with huge data for use in many different tasks.

API

A digital bridge that allows two different software to talk to each other and exchange data.

IDE

An advanced working environment where software developers can find all the tools they need while writing code.

framework

A skeleton system that makes things easier when developing software and offers ready-made structures and a set of rules.

voice cloning

A technology that analyzes a person's voice and produces artificial speech based on that person's tone of voice.
