



TreScout

DAILY TECHNOLOGY REPORT

Today's Technology Trends

August 14, 2026 (Friday)

A FEW SIPS

In today's technology agenda, new generation architectures developed for large language models and real-time video processing technologies stand out. While announcements such as Gemini 3.7 Flash and DeepSeek Harness focus on model efficiency, on the academic side, an intense work traffic on character animation in live video streams is observed.

34 HIGHLIGHTS · ~9 MIN READ

17 GitHub
Trending

5 Hacker
News

5 HuggingFace ·
Daily Models

5 HuggingFace ·
Daily Papers

2 Lobsters

01 cathrynlavery/diagram-design

Diagram-design offers 29 different editorial diagram templates designed to be used on the artificial intelligence coding assistant Claude Code. These stand-alone designs in HTML and SVG format provide a clean and simple visualization method instead of complex graphics tools.

HTML · ★ 15,284 · +4,475 today <https://github.com/cathrynlavery/diagram-design>

02 semantica-agi/semantica

Semantica is a graph-native infrastructure that provides contextual data management for artificial intelligence systems. By preserving the relational structure of the data, it ensures more reliable and auditable results in generative AI processes.

Python · ★ 6,916 · +713 today <https://github.com/semantica-agi/semantica>

03 anthropics/skills

This library, shared by Anthropic, offers reusable skill packages for artificial intelligence agents. It enables developers to standardize complex workflows and build agent-based systems faster.

Python · ★ 169,166 · +312 today <https://github.com/anthropics/skills>

04 cactus-compute/needle

Developed by Cactus Compute, Needle is a 14-megabyte foundation model that can run on small hardware such as phones, wearable devices and smart home products. This lightweight structure makes it possible to run native artificial intelligence applications on devices with limited processing power.

Python · ★ 5,082 · +769 today <https://github.com/cactus-compute/needle>

05 altic-dev/FluidVoice

FluidVoice is a dictation application that runs on macOS and performs speech-to-text conversion (STT) entirely on the device. This tool protects privacy by processing user data locally and offers an alternative solution to services such as Wispr Flow with similar functions.

Swift · ★ 9,936 · +76 today <https://github.com/altic-dev/FluidVoice>

06 unslothai/unsloth

Unsloth is a library that accelerates the training and execution of large language models (LLM) and diffusion models in a native environment. It makes fine-tuning processes of popular models more accessible by optimizing memory usage.

Python · ★ 71,157 · +328 today <https://github.com/unslothai/unsloth>

07 **macro-inc/macro**

Macro is a workspace that combines business tools such as email, chat, docs and task management into a single interface. This platform, developed with the Rust language, connects all workflows through a shared artificial intelligence memory.

Rust · ★ 2,696 · +1,239 today <https://github.com/macro-inc/macro>

08 **megadose/holehe**

Holehe is a Python tool that queries whether an entered email address is used on popular platforms such as Twitter and Instagram. It detects the account information associated with the relevant email using password reset functions.

Python · ★ 12,498 · +195 today <https://github.com/megadose/holehe>

09 **smicallef/spiderfoot**

SpiderFoot is a security tool that identifies cyber threats and maps the attack surface by automating open source intelligence (OSINT) processes. This platform, developed with Python language, analyzes the vulnerabilities of digital assets by scanning large data sources.

Python · ★ 20,712 · +283 today <https://github.com/smicallef/spiderfoot>

10 **NVIDIA-NeMo/Switchyard**

Developed by NVIDIA, Switchyard is an application interface layer that routes traffic between different providers for large language models. Thanks to its structure compatible with OpenAI and Anthropic standards, developers can flexibly switch between models according to cost and performance targets.

Rust · ★ 1,291 · +408 today <https://github.com/NVIDIA-NeMo/Switchyard>

11 **holaboss-ai/holaOS**

holaOS is an open source workspace that combines different artificial intelligence agents (Claude Code, Codex) with your tools, applications and files over a common memory. With support for more than 100 integrations and Model Link Protocol (MCP), it enables users to manage their own models (BYOK) or built-in models in a single interface.

TypeScript · ★ 6,748 · +241 today <https://github.com/holaboss-ai/holaOS>

12 **kepano/obsidian-skills**

Obsidian-skills enables artificial intelligence agents to perform operations on the note-taking application Obsidian using the command line interface (CLI) and open file formats. This tool automates personal information management by allowing agents to read and edit formats such as Markdown and JSON Canvas.

★ 45,950 · +292 today <https://github.com/kepano/obsidian-skills>

13 3b1b/manim

Developed to visualize mathematical concepts, Manim is a Python-based animation engine. It is preferred by those who produce educational content to turn complex equations and geometric structures into explanatory videos.

Python · ★ 90,959 · +176 today <https://github.com/3b1b/manim>

14 msitarzewski/agency-agents

Agency-agents is a software collection that brings together artificial intelligence agents with different areas of expertise under one roof. Performing tasks ranging from interface development to community management, these specialized agents are designed to automate specific business processes.

Shell · ★ 145,283 · +778 today <https://github.com/msitarzewski/agency-agents>

15 Lightricks/LTX-2

Published by Lightricks, LTX-2 is the official Python inference and low-order adaptation (LoRA) training suite for a generative model that produces audio and video. This tool allows developers to fine-tune the model with their own data sets.

Python · ★ 8,960 · +205 today <https://github.com/Lightricks/LTX-2>

16 lightningpixel/modly

Modly is a desktop application that creates 3D models from images and performs all operations on the local graphics card (GPU). This tool, which works without the need for an internet connection, brings the artificial intelligence-supported modeling process to personal computers.

TypeScript · ★ 5,550 · +118 today <https://github.com/lightningpixel/modly>

17 infiniflow/ragflow

RAGFlow is an open source fetch-based generation (RAG) engine that creates a context layer for large language models (LLM). It combines advanced RAG techniques with agent capabilities to optimize data processing and response quality.

Go · ★ 88,121 · +465 today <https://github.com/infiniflow/ragflow>

01 Gemini 3.7 Flash

Gemini 3.7 Flash, announced by Google, is a new generation generative artificial intelligence (generative AI) model focused on high speed and low cost. The model provides increased performance on complex reasoning tasks and offers faster response times for developers.

716 points · 393 comments <https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-3-7-flash/>

Discussion: news.ycombinator.com/item?id=49289112 (716 points · 393 comments)

02 DeepSeek Harness developer preview

DeepSeek Harness is an evaluation framework used to measure the performance of large language models. It allows developers to benchmark their models against standard test sets and systematically analyze model capabilities.

604 points · 255 comments <https://deepseek.com/harness/en/>

Discussion: news.ycombinator.com/item?id=49285244 (604 points · 255 comments)

03 Spaghettifying DRAM

Spaghettifying DRAM, which provides random number generation by taking advantage of the physical properties of the data in DRAM memory modules, works through a hardware-based vulnerability. This method increases the predictability of cryptographic keys by exploiting leaks in memory cells.

546 points · 150 comments <https://github.com/xoreaxeaxeax/skitter-creek-bath-salts>

Discussion: news.ycombinator.com/item?id=49286341 (546 points · 150 comments) · lobste.rs/s/a4zifd/unlocking_everything_on_cpu_with_dram (34 points · 8 comments)

04 Accelerating GPT-5.6 Sol Ultrafast

Cerebras, a company that produces artificial intelligence hardware, has developed a new technology that increases the inference speed for OpenAI models. This method aims to increase processing efficiency by reducing the response time of large language models.

507 points · 208 comments <https://www.cerebras.ai/blog/accelerating-gpt-5-6-sol-ultrafast-with-openai>

Discussion: news.ycombinator.com/item?id=49289844 (507 points · 208 comments)

05 Choose Boring Technology (2015)

An approach that advocates choosing proven and stable technologies in software development processes. He states that choosing highly reliable infrastructures instead of innovative tools reduces the operational burden of teams and reduces the risk of technical debt.

298 points · 147 comments <https://mcfunley.com/choose-boring-technology>

Discussion: news.ycombinator.com/item?id=49289512 (298 points · 147 comments)

01 meta-models/Muse-Glimmer-30B

Muse-Glimmer-30B, published on HuggingFace, is a multimodal language model that can process visual and textual data simultaneously. This model offers the ability to produce text-based outputs by analyzing complex images.

♥ 1,437 · 121,042 downloads <https://huggingface.co/meta-models/Muse-Glimmer-30B>

02 MiniMaxAI/MiniMax-H3

MiniMax-H3, developed by MiniMaxAI, is a generative AI model that creates videos using visual and text inputs. It allows users to produce moving images from static content or written commands.

♥ 3,844 · 1,605,940 downloads <https://huggingface.co/MiniMaxAI/MiniMax-H3>

03 Qwen/Qwen3.8-2.4T-A95B

Qwen3.8, the text-generation model developed by Alibaba, offers a high-capacity artificial intelligence architecture trained on large data sets. This system, shared on the open source model library HuggingFace, is optimized to perform complex language processing tasks.

♥ 809 · 1,012 downloads <https://huggingface.co/Qwen/Qwen3.8-2.4T-A95B>

04 Lightricks/LTX-2.5

Developed by Lightricks, LTX-2.5 is an image-to-video production model that converts still images into short video clips. This technology, used to animate visual content, allows users to create high-quality videos from static frames.

♥ 745 · 57,287 downloads <https://huggingface.co/Lightricks/LTX-2.5>

05 deepseek-ai/DeepSeek-V4-Flash-0731

DeepSeek-V4-Flash-0731, the text-generation model developed by DeepSeek, offers a structure optimized for high-performance language processing tasks. The model stands out with its fast response times among the vulnerability-heavy models shared on platforms such as Hugging Face.

♥ 3,333 · 1,431,587 downloads <https://huggingface.co/deepseek-ai/DeepSeek-V4-Flash-0731>

01 UniSwap: Streaming Audio-Visual Identity Swapping for Talking Videos

UniSwap offers a framework that simultaneously changes the visual and audio identity in speaking videos. It integrates reference visual and audio data into the video, preserving the motion, scene and timing in the source video.

1 likes · 1 comments <https://huggingface.co/papers/2608.11752>

02 LiveAnimate: Stable Long-Form Streaming Human Animation in Real-Time

LiveAnimate is a system that produces real-time human animations from a single reference image and motion stream. It optimizes the long processing times of diffusion-based models, allowing seamless video creation in interactive areas such as live broadcasts and virtual avatars.

1 likes · 1 comments <https://huggingface.co/papers/2608.11745>

03 CW-BASS v2: Saturation-Aware Pseudo-Label Selection for Semi-Supervised Segmentation under Foundation-Model Teachers

CW-BASS v2 improves semi-supervised semantic segmentation processes, which traditional methods are insufficient due to the high confidence values of foundation model teachers. This new saturation-aware approach improves model performance by optimizing the selection of labels produced by powerful models.

0 likes · 1 comments <https://huggingface.co/papers/2608.12773>

04 H2R-Bench: Benchmarking Human-to-Robot Manipulation Video Generation in World Models

H2R-Bench is a benchmarking framework developed to evaluate video production models that transfer human hand movements to robotic end-effectors. It measures the performance of world models that synthesize robot-centered manipulation videos from human videos in order to reduce the cost of data collection in robotic learning processes.

2 likes · 1 comments <https://huggingface.co/papers/2608.13049>

05 LLMRouter: Unified Infrastructure for Developing, Evaluating, and Deploying LLM Routers

LLMRouter offers a unified infrastructure to develop, evaluate and deploy routers that select the optimal large language model for different queries and budget constraints. This system breaks down the routing process into five key components, making it easy to compare the performance of different models and create scalable workflows.

48 likes · 1 comments <https://huggingface.co/papers/2608.06867>

01 I want extern "fil-c"

Software developers are looking for an external standard library to make file system operations (file I/O) in C more secure and manageable. This approach aims to reduce the margin of error in software development processes by making low-level system calls more modular.

46 points · 16 comments <https://domenkozar.com/2026/08/13/i-want-extern-fil-c/>

Discussion: lobste.rs/s/tssf5y/i_want_extern_fil_c (46 points · 16 comments)

02 python's pre-declared constants are kinda weird

Pre-declared constants in the Python language contain some inconsistencies that contradict the dynamic structure of the language. This brings into question technical limitations that make it difficult for developers to maintain full control over the code.

91 points · 15 comments <https://sebsite.pw/w/20260801-pythonconstants.html>

Discussion: lobste.rs/s/62kxco/python_s_pre_declared_constants_are_kind (91 points · 15 comments)

graph-native infrastructure

Special infrastructure designed to quickly process complex relationships between data over direct connections.

generative AI

Artificial intelligence technology that can create new texts, images or sounds based on previously learned information.

skills

Special abilities that a software or artificial intelligence has to perform a specific task.

foundation model

The basic artificial intelligence structure that is trained with very large data sets and on which different tasks can be built.

STT

Technology that listens to spoken words and converts them into written text.

LLM

Artificial intelligence models trained with huge data that can understand human language and produce text.

diffusion models

An artificial intelligence method that creates high-quality visuals from scratch by gradually editing noisy data.

fine-tuning

The process of customizing a general-purpose trained artificial intelligence with a small data set to perform better on a specific subject.

OSINT

A method of collecting and analyzing information using sources publicly available on the Internet.

large language models

Artificial intelligence models trained with huge data that can understand human language and produce text.

MCP

A protocol that enables artificial intelligence models to connect to different data sources and tools in a standard way.

BYOK

A security approach that allows the user to encrypt their data using their own security keys.

CLI

Text-based interface that runs programs with written commands rather than the mouse.

generative model

A system that can produce new and original content using patterns in training data.

inference

The process of a trained artificial intelligence model producing a result by processing a new input given to it.

LoRA

A method of training a large artificial intelligence model more quickly and efficiently by updating only a small part of it, without changing the entire model.

TreScout · trescout.com · Every day, on your schedule.