



TreScout

DAILY TECHNOLOGY REPORT

Today's Technology Trends

September 26, 2026 (Saturday)

A FEW SIPS

Today's technology agenda is defined by the security vulnerabilities of AI agents and open-source decision models. Discussions on Hacker News and GitHub focus on the risky interactions of autonomous systems on platforms like Hugging Face and the concept of commodified intelligence.

34 HIGHLIGHTS · ~8 MIN READ

16 GitHub
Trending

5 Hacker
News

5 HuggingFace ·
Daily Models

5 HuggingFace ·
Daily Papers

3 Lobsters

Data captured September 26, 2026 at 12:07 PM.

Report source counts and item metadata are fixed at capture time. Discovery pages may show newer live source values.

01 paperclipai/paperclip

Paperclip is an open-source application that allows you to manage AI agents in the workplace from a single hub. Developed with TypeScript, this tool makes it easier to coordinate the workflows of different agents and standardize operational processes.

TypeScript · ★ 85,663 · +2,109 today <https://github.com/paperclipai/paperclip>

02 anthropics/claude-plugins-official

Anthropic gathers high-quality plugins developed for its AI assistant, Claude, in a single hub. The Claude Code Plugins directory offers Python-based tools that developers can integrate into their projects in a standardized structure.

Python · ★ 37,014 · +83 today <https://github.com/anthropics/claude-plugins-official>

03 vectorize-io/hindsight

Hindsight is a Python library that enables AI agents to create long-term memory by learning from past interactions. It allows developers to give their agents experiential learning capabilities, supporting the generation of more consistent and personalized responses over time.

Python · ★ 30,381 · +1,653 today <https://github.com/vectorize-io/hindsight>

04 obra/superpowers

Superpowers, developed by Obra, offers an agentic skills framework and software development methodology for AI agents. This tool aims to standardize the workflows of autonomous systems that perform complex tasks.

Shell · ★ 291,794 · +468 today <https://github.com/obra/superpowers>

05 mattpocock/skills

This repository shared by Matt Pocock offers ready-to-use skill packages that software developers can use for AI agents. This structure, which developers can use in their own projects, aims to make agent-based workflows more efficient.

Shell · ★ 269,905 · +583 today <https://github.com/mattpocock/skills>

06 dream-num/univer

Univer is an open-source platform that combines office tools such as spreadsheets, documents, and presentations into a single runtime. Developed to facilitate the interaction of AI agents with office software, this structure allows for managing different data types in a single interface.

TypeScript · ★ 18,873 · +1,050 today <https://github.com/dream-num/univer>

07 anthropics/skills

This repository shared by Anthropic offers reusable skill packages for AI agents. It creates a Python-based library structure that allows developers to standardize complex workflows.

Python · ★ 178,454 · +189 today <https://github.com/anthropics/skills>

08 androoAGI/starnet

Starnet is a desktop management tool that enables AI agents to perform real tasks on a computer with a local-first working principle. It allows users to define their own keys and monitor the workflows of agent teams through a pixel art interface.

JavaScript · ★ 575 · +93 today <https://github.com/androoAGI/starnet>

09 derv82/wifit3

Wifit3 is a Python tool developed to test wireless network security, which works only with USB devices and supports different operating systems. It is a cross-platform compatible alternative to Wifite, the popular security software that automates network auditing processes.

Python · ★ 1,060 · +183 today <https://github.com/derv82/wifit3>

10 kelseyhightower/kubernetes-the-hard-way

Kubernetes the Hard Way, prepared by Kelsey Hightower, is a guide that teaches how to install the container management platform Kubernetes from scratch without using any automation scripts. This work provides a technical training resource for users who want to understand the system's core components and operating principles through manual configuration processes.

★ 50,174 · +119 today <https://github.com/kelseyhightower/kubernetes-the-hard-way>

11 rohitg00/ai-engineering-from-scratch

ai-engineering-from-scratch, which teaches software developers the processes of building AI applications from the ground up, offers a Python-based training guide. The project explains the entire development cycle, from training AI models to deploying them, with practical examples.

Python · ★ 57,797 · +1,177 today <https://github.com/rohitg00/ai-engineering-from-scratch>

12 shy3130/tick-stock-panel

Tick-stock-panel is a customizable stock selection, monitoring, and backtesting platform that works with Chinese stock market (A-share) data. It is an open-source Python tool that automates strategy development and stock analysis processes with the support of large language models.

Python · ★ 5,204 · +44 today <https://github.com/shy3130/tick-stock-panel>

13 google/ax

Ax, developed by Google, offers an open-source orchestration runtime for AI agents. Written in the Go language, this platform facilitates the management and execution of complex agentic workflows.

Go · ★ 11,703 · +1,379 today <https://github.com/google/ax>

14 **NVIDIA/Model-Optimizer**

Model Optimizer, developed by NVIDIA, provides performance improvements by compressing deep learning models using methods such as quantization, distillation, and pruning. This library prepares optimized models to run faster on inference engines like TensorRT-LLM.

Python · ★ 4,563 · +359 today <https://github.com/NVIDIA/Model-Optimizer>

15 **pbakaus/impeccable**

Impeccable is a design language developed to improve the design outputs of AI models. This JavaScript-based tool aims to make the visual and structural design decisions of AI systems more consistent.

JavaScript · ★ 71,352 · +306 today <https://github.com/pbakaus/impeccable>

16 **openbao/openbao**

OpenBao is security software developed to manage, store, and distribute sensitive data such as passwords, certificates, and encryption keys. Written in the Go language, this tool provides a central control point to ensure the security of critical data.

Go · ★ 7,846 · +49 today <https://github.com/openbao/openbao>

01 Revealing the details of how OpenAI agents hacked Hugging Face

Swarmtraces provides an analysis detailing how AI agents exploit a vulnerability on Hugging Face. The study examines the methods and security vulnerabilities used by autonomous agents when gaining access to systems from a technical perspective.

468 points · 283 comments <https://swarmtraces.org/>

Discussion: news.ycombinator.com/item?id=49849985 (468 points · 283 comments)

02 Ollaya – Ollama for open-source, Jev-style decision models

Ollaya is a platform that enables running decision models on Ollama, a tool for running local language models. It allows users to automate complex logical processes with open-source models.

443 points · 117 comments <https://ollaya.dev/>

Discussion: news.ycombinator.com/item?id=49848269 (443 points · 117 comments)

03 U.S. appeals court upholds designation of Anthropic as supply chain risk

A US appeals court has upheld a decision that AI company Anthropic carries supply chain risks. This ruling creates restrictions on Anthropic's participation in federal tenders and its status in public projects, as the company develops defense technologies.

441 points · 762 comments <https://www.cnn.com/2026/09/25/pentagon-anthropic-ai-risk-appeals-court.html>

Discussion: news.ycombinator.com/item?id=49845977 (441 points · 762 comments)

04 Plan mode is dead

In AI-powered coding tools, the "plan mode," which requires planning from the user beforehand, is being replaced by direct implementation phases. Developers prefer faster workflows where the AI writes and corrects code in real-time instead of preparatory stages that break down complex processes.

307 points · 286 comments <https://www.aymannadeem.com/artificial/intelligence,/developer/tools/2026/09/24/pl-an-mode-is-dead.html>

Discussion: news.ycombinator.com/item?id=49840054 (307 points · 286 comments)

05 First Principles Thinking

First principles thinking is an analytical method that allows for rebuilding complex problems by breaking them down into their smallest, verifiable parts. Instead of copying existing assumptions or traditional methods, this approach focuses on generating original solutions by analyzing the fundamental facts of a subject.

259 points · 112 comments <https://sunilsadasivan.com/writing/first-principles-thinking/>

Discussion: news.ycombinator.com/item?id=49844736 (259 points · 112 comments)

01 convaiinnovations/laya

Laya, shared on HuggingFace, is an AI model developed for text classification tasks. This tool, which allows users to sort text data into specific categories, offers a practical solution in natural language processing workflows.

♥ 3,752 · 0 downloads <https://huggingface.co/convaiinnovations/laya>

02 Qwen/Qwen-Image-2.1

Qwen-Image-2.1, developed by Alibaba, is a text-to-image model that generates visuals from text prompts. Shared on the open-source AI platform HuggingFace, this model offers a high-performance alternative in visual production processes.

♥ 2,353 · 42,469 downloads <https://huggingface.co/Qwen/Qwen-Image-2.1>

03 abenzerps/Qwen-Image-2.1-Uncensored-GGUF

Qwen-Image-2.1-Uncensored, shared by Abenzerps, is an artificial intelligence model capable of text-to-image generation. With its uncensored structure, the model offers users a wider range of freedom in their content creation processes.

♥ 1,847 · 715,906 downloads <https://huggingface.co/abenzerps/Qwen-Image-2.1-Uncensored-GGUF>

04 XingChen-AGI/Xing4.0-29B-A4B

Xing4.0-29B-A4B, developed by XingChen-AGI, is a large language model focused on text generation. Shared on HuggingFace, this model is used to perform various natural language processing tasks with its large parameter structure.

♥ 1,692 · 42,950 downloads <https://huggingface.co/XingChen-AGI/Xing4.0-29B-A4B>

05 prism-ml/Ternary-Bonsai-2-27B-gguf

Ternary-Bonsai-2-27B, developed by Prism-ml, is an artificial intelligence model optimized for text generation. Offered in GGUF format, this model aims to run high-performance language processing tasks with lower hardware requirements.

♥ 2,105 · 3,109,078 downloads <https://huggingface.co/prism-ml/Ternary-Bonsai-2-27B-gguf>

01 Learning to Discover Interesting Mathematics

Large language models are gaining proficiency in solving complex mathematical problems and generating new theorems. Researchers are defining a new method based on the relationship between theorem proof length and information content to measure the original value and utility of the mathematical knowledge produced by these models.

6 likes · 2 comments <https://huggingface.co/papers/2609.28603>

02 RGBD20K: A Large-Scale Benchmark for RGB-D Semantic Segmentation

RGBD20K is a large-scale dataset developed for training semantic segmentation models that combine color image and depth data (RGB-D). Covering 160 different categories, this resource offers higher diversity and detailed labeling compared to existing popular datasets (NYUv2, SUN RGB-D).

7 likes · 2 comments <https://huggingface.co/papers/2609.29028>

03 AV-GRPO: Modality-Anchored Decoupling Diffusion Reinforcement Learning for Joint Audio-Video Generation

AV-GRPO is a modality-focused reinforcement learning method developed to solve coordination problems in audio and video generation. By separating reward signals from different data types, it aims to improve synchronization between audio and video and their alignment with text.

3 likes · 2 comments <https://huggingface.co/papers/2609.29816>

04 Your Transformer Can Hold Two Thoughts at Once: Evidence of Linear Superposition in LLMs

The Transformer architecture underlying large language models (LLMs) can simultaneously provide both probability distributions as output when inputs from different text streams are combined. This phenomenon, called the Superposition Linearity Hypothesis, shows that models exhibit a fundamental linearity despite their complex structures.

58 likes · 2 comments <https://huggingface.co/papers/2609.29845>

05 Just Ask Jev: Reinforcement Learning for Calibrated Decisions as a Zero-Shot Detector of AI Alignment Failures

Jev is a model trained with the Reinforcement Learning with Calibrated Decisions (RLCD) method that provides probabilistic answers to multiple questions about text inputs in a single call. Unlike existing classifiers used to detect alignment failures in language models, it offers a more flexible and efficient evaluation process.

2 likes · 2 comments <https://huggingface.co/papers/2609.29429>

01 Commodified Intelligence

The commoditization of artificial intelligence models allows intelligence to be positioned as a standard input in software development processes. This enables complex cognitive capabilities to be consumed as easily accessible and scalable commodity services through application development platforms.

29 points · 27 comments <https://herecomesthemoon.net/2026/09/commodified-intelligence/>

Discussion: lobste.rs/s/nidcls/commodified_intelligence (29 points · 27 comments)

02 What About Rails?

Ruby on Rails is being re-evaluated in modern web development processes for its speed and efficiency-focused structure. Software developers are choosing the monolithic framework offered by Rails as an alternative to complex microservice architectures.

46 points · 16 comments <https://jardo.dev/what-about-rails>

Discussion: lobste.rs/s/rvvqos/what_about_rails (46 points · 16 comments)

03 Goodbye Google

Software developer Robert O'Callahan explains that he has stopped using Google's services due to the decline in quality in Google's search results and its ad-focused strategies. He is restructuring his personal digital ecosystem by switching to independent alternatives instead of the Google search engine.

84 points · 16 comments <https://robert.ocallahan.org/2026/09/goodbye-google.html>

Discussion: lobste.rs/s/sxlf4a/goodbye_google (84 points · 16 comments)

AI agents

Software that can make decisions on its own and perform specific tasks without human intervention.

plugins

Small add-ons used to add new capabilities to the existing features of a software.

long-term memory

The ability of artificial intelligence to remember past information and user preferences for a long time and use them in subsequent operations.

agentic skills framework

A structural system that organizes the capabilities software needs to solve complex tasks step by step.

skills

Specialized skills that a software possesses to perform a specific task.

runtime

The duration during which a program is executed on a computer and actively performs its task.

local-first

The principle of storing and processing data directly in the device's own memory without needing the internet.

A-share

A type of stock traded on Chinese stock exchanges and reserved for local investors.

backtesting

Testing how successful an investment strategy would be using historical data.

large language models

Artificial intelligence models trained on massive amounts of data that can understand human language and generate text.

quantization

Saving data with lower precision to reduce the space occupied by artificial intelligence models.

distillation

The process of transferring the knowledge of a large and complex artificial intelligence model to another smaller and faster-running model.

pruning

Lightening the model by clearing unnecessary connections that do not affect the performance of the artificial intelligence model.

inference engine

The software component that enables a trained artificial intelligence model to make predictions on new data.

design language

A set of consistent design rules that determine the visual style, colors, and usage patterns of a product.

autonomous agents

Intelligent software that creates its own strategy and acts independently to achieve defined goals.

TreScout · trescout.com · Every day, on your schedule.