



TreScout

FRESH ONLY · DAILY TECHNOLOGY REPORT

Today's New Trends

September 25, 2026 (Friday)

A FEW SIPS

This report contains trends that have emerged for the first time today, without repeating records included in normal reports from the last 30 days.

18 HIGHLIGHTS · ~5 MIN READ

5 GitHub Trending

5 Hacker News

5 HuggingFace · Daily Papers

3 Lobsters

Data captured September 25, 2026 at 12:24 PM.

Report source counts and item metadata are fixed at capture time. Discovery pages may show newer live source values.

01 vectorize-io/hindsight

Hindsight offers a learning memory layer for AI agents. This open-source library, which improves the decision-making processes of agents by inferring from past interactions, allows systems to produce more consistent results over time.

Python · ★ 28,397 · +1,668 today <https://github.com/vectorize-io/hindsight>

02 NVIDIA/Model-Optimizer

NVIDIA Model Optimizer is a library that compresses deep learning models using techniques such as quantization, distillation, and pruning. This tool prepares models for use in inference engines like TensorRT-LLM to increase inference speed.

Python · ★ 4,209 · +44 today <https://github.com/NVIDIA/Model-Optimizer>

03 FxEmbed/FxEmbed

FxEmbed is a content embedding tool that ensures posts from X and Bluesky platforms appear correctly in applications like Discord or Telegram. It ensures that media elements such as multiple images, videos, and polls are displayed completely on messaging platforms.

TypeScript · ★ 5,416 · +182 today <https://github.com/FxEmbed/FxEmbed>

04 julyx10/lap

Lap, developed with Vue, is an offline-first photo manager that allows you to manage large local photo archives without an internet connection. It enables users to organize their personal libraries quickly and securely via the browser.

Vue · ★ 2,989 · +122 today <https://github.com/julyx10/lap>

05 leejet/stable-diffusion.cpp

Stable-diffusion.cpp is an inference library that runs image generation models (such as Stable Diffusion, Flux, and Wan) in pure C/C++. This tool allows high-performance models to be run on low-end systems or specialized environments without additional dependencies.

C++ · ★ 7,316 · +36 today <https://github.com/leejet/stable-diffusion.cpp>

01 F-Droid 2.0

F-Droid, the free software repository for Android, is modernizing the app discovery experience with its 2.0 version by revamping its user interface. The update brings performance improvements and new design standards to make the open-source software ecosystem more accessible.

1,232 points · 344 comments <https://f-droid.org/2026/09/24/f-droid-2.0-a-new-chapter-for-android-freedom.html>

Discussion: news.ycombinator.com/item?id=49831968 (1,232 points · 344 comments)

02 Show HN: Make cursed fonts like Times New Bastard

Times New Bastard is an entertainment-focused tool that blends different fonts within a single text by using the context-based ligature feature of the OpenType font format. It offers a fast conversion process by running Python code on the client side using WebAssembly (WASM) technology running in the browser.

686 points · 93 comments <https://bastardica.mitpit.com/>

Discussion: news.ycombinator.com/item?id=49823738 (686 points · 93 comments)

03 Two-tier encryption in the UK

The UK government is working on a two-stage encryption model that would provide access to message content on platforms using end-to-end encryption. This regulation aims to grant law enforcement legal oversight authority while maintaining privacy standards.

447 points · 398 comments <https://macanorak.com/two-tier-encryption-in-the-uk/>

Discussion: news.ycombinator.com/item?id=49828731 (447 points · 398 comments)

04 Why is the liver so weirdly regenerative?

The liver's high regenerative capacity is an evolutionary defense mechanism against the damage the organ is exposed to while performing its duty of clearing toxins from the body. This biological feature explains the liver's unique functioning in terms of cell division and tissue repair.

403 points · 206 comments <https://dynamight.substack.com/p/liver>

Discussion: news.ycombinator.com/item?id=49832938 (403 points · 206 comments)

05 Rails World 2026 Opening Keynote [video]

The opening keynote of the Rails World 2026 conference, held for Ruby on Rails developers, shares the future vision and new technical features of the Ruby on Rails web framework. The presentation includes information about updates focused on improving the developer experience and the platform's long-term roadmap.

333 points · 365 comments https://www.youtube.com/watch?v=vDjW_dRyKXY

Discussion: news.ycombinator.com/item?id=49817680 (333 points · 365 comments)

01 Parts-of-Speech as Emergent Categories in SAE Latent Space

Sparse autoencoders are a method used to examine the internal representations of language models. Researchers reveal that parts-of-speech categories are encoded in the hidden layers of these models as structured groups rather than individual features.

6 likes · 1 comments <https://huggingface.co/papers/2609.29362>

02 Coding Agents for Generalized Task and Motion Planning Problems

Coding agents are used to automate geometric and dynamic constraints in task and motion planning (TAMP) problems. This approach leverages code synthesis capabilities to reduce planning effort in new problem instances.

3 likes · 1 comments <https://huggingface.co/papers/2609.30233>

03 DeltaWAM: Delta World Action Models for Bimanual Manipulation

DeltaWAM was developed to address inefficiencies in world-action models that enable robots to perform bimanual manipulation. Instead of processing all video frames, the system accelerates robotic control processes by modeling only the changes (deltas).

1 likes · 1 comments <https://huggingface.co/papers/2609.28811>

04 Rufus-Air: An Open LLM Post-Training Recipe

Rufus-Air is an eight-stage post-training method applied to the large language model GLM-4.5-Air-Base. This process offers a progressive reinforcement learning pipeline that extends from basic capabilities to complex reasoning and agent skills.

1 likes · 1 comments <https://huggingface.co/papers/2609.29421>

05 Neural Spectral Capacity: Measuring and Designing Architectures from Network Specification Alone

Neural Spectral Capacity is a new method that measures the architectural structure of artificial intelligence models based on the singular value spectrum of weight matrices. Unlike traditional metrics such as parameter count and processing power (FLOPs), it aims to evaluate the model's internal architectural efficiency more precisely.

2 likes · 1 comments <https://huggingface.co/papers/2609.23087>

01 **Breaking Up with Google Play: Why Conversations Is Now Free**

The XMPP-based messaging application Conversations has become completely free by ending its reliance on payment systems in the Google Play store. The developer encourages users to obtain the application directly from their own website or through the open-source app store F-Droid.

116 points · 16 comments <https://guiltsch.de/posts/breaking-up-with-google-play/>

Discussion: lobste.rs/s/4jqfz1/breaking_up_with_google_play_why (116 points · 16 comments)

02 **SourceHut account takeover via build logs (XSS in ansi2html.py)**

A security vulnerability found in the build logs of the software development platform SourceHut allows for account takeover via cross-site scripting (XSS). The issue stems from faulty character handling logic within the text formatting tool ansi2html.py.

66 points · 3 comments <https://blog.arusekk.pl/posts/srht-account-takeover/>

Discussion: lobste.rs/s/kylcr0/sourcehut_account_takeover_via_build (66 points · 3 comments)

03 **I asked Meta's Muse for its filesystem and it sent me 6.8 GB**

The artificial intelligence model Muse, developed by Meta, sent 6.8 GB of data in response to a user's file system request. This situation highlights the risks regarding access to runtime environments and data security in large language models.

58 points · 10 comments <https://mouse.dev/blog/muse-runtime-export/>

Discussion: lobste.rs/s/mdczxv/i_asked_meta_s_muse_for_its_filesystem_it (58 points · 10 comments)

memory layer

A digital memory layer that enables artificial intelligence models to remember past interactions and store information.

runtime

The moment software is executed on a computer and the environment it requires during this process.

quantization

A method of lightening an AI model by reducing the precision of numerical values to decrease the space it occupies.

distillation

The process of transferring the knowledge of a large and complex AI model to a smaller and faster model.

pruning

The process of making an AI model more efficient by deleting unnecessary connections that do not affect its performance.

inference engines

A software engine that enables a trained AI model to make predictions on new data.

embed

The process of converting data into numerical lists so that computers can understand the semantic relationships between them.

large language models

AI models trained on massive amounts of data that can understand human language and generate text.

framework

A foundational structure that facilitates software development by providing a set of ready-made tools and rules.
